## **Project 3: Unix Signals and File I/O**

**Due**: Monday 29 April 2013, 23:59:00 (11:59 PM) Pacific USA time zone.

Points on this assignment: 150 points with 10 bonus points available.

Work submitted late will be penalized as described in the course syllabus. You must submit your work twice for this and all other homework assignments in this class. Ecampus wants to archive your work through BlackBoard and EECS needs you to submit through TEACH to be graded. If you do not submit your assignment through TEACH, it cannot be graded (and you will be disappointed with your grade). Make sure you submit your work through TEACH. Submit your work for this assignment as a single bzip file through TEACH. The same single bzip file should also be submitted through BlackBoard.

\_\_\_\_\_

Place all of the files you produce for this assignment in a single directory, maybe something called Homework3.

In this assignment, you will be working with Unix signals and file I/O system calls. This assignment is a significant step up in complexity from the previous 2 homework assignments. This assignment will also be only in C, with some Makefile fun mixed in. If you look through some of the C code on the previous 2 homeworks, you can find some C code fragments that may be helpful in this assignment. You have 2 weeks to complete this assignment. However, I urge you to not delay beginning it.

- 1. **10 points**. When you are ready to submit your files for this assignment, make sure you submit a <u>single bzip file</u>. Review homework #1, problem #1 if you need a refresher on how to do this. If your file is not a <u>single bzip file</u>, you cannot receive points on this assignment.
- 2. **20 points**. Write a C program to send and receive Unix signals. Don't go overboard on this portion of the assignment. You should be able to easily complete this portion with less than 50 lines of C code. It may take you less time to write the C code than it took me to write the problem description.

You will write an application in C (sig\_demo.c) that installs 3 different functions as signal handlers. Your C code will handle the following signals: SIGUSR1, SIGUSR2, and SIGINT. It may be tempting to use a single signal handler for all 3 signals, but don't. Create a separate single handler function for each signal.

When your application receives the SIGUSR1 signal, it should print:

Caught SIGUSR1

Your application should not exit when the SIGUSR1 signal is handled.

When your application receives the SIGUSR2 signal, it should print:

```
Caught SIGUSR2
```

Your application should not exit when the SIGUSR2 signal is handled.

When your application receives the SIGINT signal, it should print:

```
Caught SIGINT, exiting
```

When it receives the SIGINT signal, your application **should** exit (after printing the message of course).

You may call printf() for these messages within your signal handler. I know that it is not strictly *safe* to call printf() within a signal handler, but we will this time.

Once your application has installed the 3 signal handlers, it should send the signals to itself, in this order: SIGUSR1, SIGUSR2, SIGINT. The getpid() system call can be your friend here.

When I run your compiled program, I should see the following

```
$ ./sig_demo
Caught SIGUSR1
Caught SIGUSR2
Caught SIGINT, exiting
$
```

3. **90 points total**. This portion of the homework is about reading and writing files. You will need to stat() or fstat() files, check file permissions, check file time stamps, and perform a seek() into a file.

Write a C program on os-class or CS311 VM called myar.c. This program will illustrate the use of file I/O on Unix by maintaining a Unix archive library, in the standard archive format.

Once compiled your program should run in a manner similar to the standard Unix command ar. You can look at the man page for ar for insight. However, for this assignment, the following is the syntax your program must support:

```
myar key archive-file [member [...]]
```

where archive-file is the name of the archive file to be used, and key is one of the following options:

-d	Quickly append named files (members) to archive. 15 points (of the 90).
	Check the meaning of append in the notes below.
-X	Extract named members. 15 points (of the 90). Just as the regular ar
	command, if no member is named on the command line when extracting
	files, all files are extracted from the archive. The permissions on the
	extracted files should match permissions on the files before archiving
	(as described in the notes below).
-t	Print a concise table of contents of the archive. 5 points (of the 90).
	The concise table of contents for your application (myar) should match
	exactly the output from the "ar t" command on the same archive file.
-A	Print a verbose table of contents of the archive. 5 points (of the 90). The
	verbose table of contents for your application (myar) should match
	exactly the output from the "ar tv" command on the same archive
	file. See the man page on ar.
-d	Delete named files from archive. 20 points (of the 90). Make sure you
	read the note below about the -d option and creation of a new file.
-A	Quickly append all "regular" files in the current directory (Except the
	archive itself). 20 points (of the 90). There is not an option for the
	Unix ar command that does this.
-M	Extra credit 5 points: for a given timeout (in seconds), add all
	<u>modified</u> files to the archive (Except the archive itself). <b>See note h.</b>
	There is not an option for the Unix ar command that does this.

The archive file maintained must use <u>exactly</u> the standard format defined in /usr/include/ar.h, and in fact may be tested with archives created with the ar <u>command</u>. Do not copy or in any way modify the ar.h include file.

The options listed above are compatible with the options having the same name in the ar command, except for the following exceptions: the -v and -t command take no further argument, and list all files in the archive. -v is short for -t -v on the ar command. The -A and -w commands are new options not in the usual ar command.

## Notes:

- a) For the -q command myar should create an archive file if it doesn't exist, using permissions "666". For the other commands myar reports an error if the archive does not exist, or is in the wrong format.
- b) You will have to use the system calls stat and utime to properly deal with extracting and restoring the proper timestamps. Since the archive format only allows one timestamp, store the mtime and use it to restore both the atime and mtime. Permissions should also be restored to the original value, subject to umask limitation.
- c) The -q and -A commands do not check to see if a file by the chosen name already exists. It simply appends the files to the end of the archive.

- d) The -x and -d commands operate on the first file matched in the archive, without checking for further matches. It is possible for a file name to exist more than once in an archive, use the first one that matches.
- e) In the case of the -d option, you will have to build a new archive file to recover the space. Do this by unlinking the original file after it is opened (or after you've completed reading it), and creating a new archive with the original name.
- f) Since file I/O is expensive, do not make more than one pass through the archive file, an issue especially relevant to the multiple delete case.
- g) You are required to handle multiple file names as members.
- h) For the -w flag, the command will take as long as specified by the timeout argument. You should print out a status message upon adding a new file. This may result in many different copies of the same file in the archive.
- i) Make sure you lookup what a "regular" file is in Unix.
- j) You must have a Makefile for this portion of the homework. Your Makefile must contain the following 2 targets:
  - 1. all -- the all target should compile and link the final program. The all target should also be the default target (which only means that it is the first target in the Makefile).
  - 2. clean the clean target should remove all binary applications (myar) and object files (.o's). In addition, you should remove any editor droppings from emacs, vi, or whatever. The sample Makefile from Homework1 includes a clean target. Make sure you use this before you bundle all your files together for submission.

The Makefile may (should) contain other targets the build any out of date modules. Again, the sample Makefile from Homework1 gives some guidance on this. Targets like all and clean are sometimes called phony targets because they don't actually produce a file.

When I build your program, I should be able to just type make to have it completely build.

- k) For **5 points extra credit**, any time a file is added that already exists, remove the old copy from the archive, but <u>only</u> if it is not the same. If identical, do not add the new file. Make sure you clearly comment your meaning of "identical."
- l) It is not necessary that you use getopt() to process argy from the command line. If you have a simpler method to use, that is fine.
- m) I have created some sample files that you can use for testing your application. They are all plain text, so you can actually just cat the archive file after you've create it to see how it looks and compare it to one created with ar. The sample files are: 1-s.txt, 2-s.txt, 3-s.txt, 4-s.txt, and 5-s.txt. One of the things you'll note about the test files that some of them have an even number of bytes and others have an odd number of bytes. This will be important to you.

- page/opensolaris/3head/ar.h/ One of the other things you learn in this assignment is how to carefully read specifications and man pages. Just take my word for this and carefully read the 2 web pages.
- o) Since there are quite a few parts to completing this problem in the homework, I suggest that you start small. Start with just creating code that performs the -t option on an archive file you've created with ar. From there, move to the -v option. From there, I'd go -q, -x, -d, -A, and -w, but go in the order that make a clear progression for you.
- p) The length of file names used will not exceed 15 characters for this assignment.
- q) Make sure your code compiles before you submit it. Simply being able to compile your code is worth 5 points (of the 90).
- r) Because you are potentially working with binary files, you should use the read() and write() system calls, not fscanf() and fprintf(), for the file I/O. You can use printf() for terminal output.
- 4. **30 points**. Put tests as targets in your Makefile. Since you need to have a number of tests for your code, put them into the Makefile so it is easy for you to run then easily and consistently. This is a simple form of regression testing.
  - a. Put a target called testq12345 in your Makefile. The testq12345 target will do the following:
    - i. Remove any file named ar12345.ar. The file ar12345.ar does not exist, don't show an error.
    - ii. Remove any file named myar12345.ar. The file myar12345.ar does not exist, don't show an error.
    - iii. Create a file named ar12345.ar using this call:

```
ar q ar12345.ar 1-s.txt 2-s.txt 3-s.txt 4-s.txt 5-s.txt
```

iv. Create a file named ar12345.ar using this call:

```
myar -q myar12345.ar 1-s.txt 2-s.txt 3-s.txt 4-s.txt 5-
s.txt
```

v. Compare the files create by ar and myar.

```
diff ar12345.ar myar12345.ar
```

- vi. The result of the diff command should show no differences.
- b. Put targets in your Makefile called testq135 and testq24 which are like the target testq12345, but only use the noted subset of files (1, 3, 5, 5)and 2, 4).
- c. Put a target in your Makefile called testq that will run the testq12345, testq135, and testq24 targets. 5 Points (of the 30).
- d. Put a target called testt12345 in your Makefile. The testt12345 target will do the following:
  - i. Remove any file named ar12345.ar. The file ar12345.ar does not exist, don't show an error.
  - ii. Create a file named ar12345.ar using this call: ar q ar12345.ar 1-s.txt 2-s.txt 3-s.txt 4-s.txt 5-s.txt
  - iii. Run the following commands:

```
ar t ar12345.ar > ar-ctoc.txt
myar -t ar12345.ar > myar-ctoc.txt
```

iv. Compare the table of contents files by using this command:

```
diff ar-ctoc.txt myar-ctoc.txt
```

- v. The result of the diff command should show no differences.
- e. Put targets in your Makefile called testt135 and testt24 which are like the target testt12345, but only use the noted subset of files (1, 3, 5, and 2, 4).
- f. Put a target in your Makefile called testt that will run the testt12345, testt135, and testt24 targets. 5 Points (of the 30).
- g. Put a target called testv12345 in your Makefile. The testv12345 target will do the following:
  - i. Remove any file named ar12345.ar. The file ar12345.ar does not exist, don't show an error.
  - ii. Create a file named ar12345.ar using this call:
    ar q ar12345.ar 1-s.txt 2-s.txt 3-s.txt 4-s.txt 5-s.txt
  - iii. Run the following commands:
     ar tv ar12345.ar > ar-vtoc.txt
     myar -v ar12345.ar > myar-vtoc.txt
  - iv. Compare the table of contents files by using this command: diff ar-vtoc.txt myar-vtoc.txt
  - v. The result of the diff command should show no differences.
- h. Put targets in your Makefile called testv135 and testv24 which are like the target testv12345, but only use the noted subset of files (1, 3, 5, and 2, 4).
- i. Put a target in your Makefile called testv that will run the testv12345, testv135, and testv24 targets. 5 Points (of the 30).
- j. Put a target in your Makefile called tests that will run the testq, testt, and testv targets. With this, you should only need to run "make tests" to run a fairly complete set of tests on your code. Being able to easily run a consistent set of tests on your code and help you quickly locate inadvertent changes in your program that might not have been caught until much later. Think about using macros in your Makefile for some of these. Test early and test often. 15 Points (of the 30).
- k. I'm sure you will want and need to have additional tests; you can put them in your Makefile as well.

\_\_\_\_\_\_

Things to include with the assignment (in a single bzipped file):

- 1. C source code for the solutions to the posed problem (all files).
- 2. A Makefile to build your code.
- 3. A file describing how you ran your tests.

A file describing how you ran your tests, including where you found input files Please combine all of the above files into a single bzip file prior to submission.